

Mayo Clinic: Building the Platform for the Future of Medicine

ITEC-617 | Kogod School of Business, American University | Spring 2026

This case was developed with AI assistance using the Agentic AI Case Study Development methodology. All quotes and data points are sourced from publicly available materials verified against original sources. Intended for classroom discussion only. Not an endorsement or illustration of effective or ineffective management.

The Question on the Table

In the spring of 2020, Dr. John Halamka sat in his new office in Rochester, Minnesota, with a question that would define the next decade of American healthcare: could a 157-year-old medical institution remake itself as a technology platform -- and should it?

Three months earlier, Halamka had been at Harvard Medical School, teaching what he knew about health IT to a generation of students who had grown up with smartphones but were now entering a system still faxing patient records. He had been offered something different: the chance to build an entirely new kind of healthcare organization from inside one of the most trusted brands in medicine. He said yes.

Halamka had arrived at Mayo Clinic on January 1, 2020 as the inaugural President of Mayo Clinic Platform, a newly created business unit charged with turning the institution's most valuable -- and most sensitive -- asset into a competitive advantage. That asset was data: decades of de-identified patient records representing conditions, treatments, and outcomes at a scale no algorithm had ever accessed.

Then, in March 2020, COVID-19 arrived. And suddenly, the question of what a hospital could do with its data became the most urgent question in global health.

"When I look at the 40 years that I've been in this industry ... I've often looked at 'when is there a perfect storm for innovation?'" -- Dr. John Halamka, HealthLeaders Media, 2020

He believed he had found one.

Mayo Clinic: Institution, Not Just Hospital

To understand what Halamka was attempting, you first have to understand what Mayo Clinic is -- and what it is not.

Mayo Clinic is not a hospital chain in the commercial sense. It is a nonprofit integrated medical practice founded in 1863 in Rochester, Minnesota, organized around a mission it has never changed: to inspire hope and contribute to health and well-being by providing the best care to every patient through integrated clinical practice, education, and research. The clinic grew from a small surgical practice started by Dr. William Worrall Mayo into the most recognized name in American medicine, built on a

philosophy of group practice and physician collaboration at a time when individual practitioners dominated the field.

By 2023, that mission had produced a \$17.9 billion organization -- revenues up 10.2% from \$16.3 billion in 2022 -- with an operating income of \$1.1 billion and a 6% operating margin. Mayo participated in 2,500 clinical trials and served approximately 1.3 million patients from more than 130 countries in that year alone. It employed approximately 86,000 people across campuses in Minnesota, Arizona, and Florida. (Fierce Healthcare, February 2024)

But the financial figures obscured what made Mayo genuinely distinct. Mayo patients did not typically come for routine care. They came because other physicians had run out of answers: the rare cancer whose histology no one recognized, the autoimmune condition that had stumped six specialists, the surgical challenge that required a team of twenty working in synchronized concert. Mayo had spent generations becoming the institution of last resort -- and first resort -- for complexity.

That diagnostic depth generated data unlike anything in commercial healthcare. Not just claims data or routine encounter records, but longitudinal patient histories tracked across decades, specialties, and outcomes. Genomic sequences paired with 30-year clinical trajectories. Imaging studies correlated with pathology and long-term survival. The clinical notes of 35,000 physicians, written across four decades of complex patient presentations, in unstructured language that captured the texture of medical reasoning in a way that structured fields never could.

Other health systems had more patients. None had this depth of complexity. And it was precisely the complexity -- the rare diseases, the edge cases, the puzzles -- that made Mayo's data most valuable for training AI that needed to learn from the hardest problems, not just the most common ones.

In 2019, Mayo's leadership made a decision: this data could be the foundation of something entirely new. The question was how to monetize it without betraying the patients who had generated it.

John Halamka: The Right Person for an Impossible Job

Dr. John Halamka was, by late 2019, one of the most recognizable figures in health information technology -- and one of the most unusual. He was simultaneously a practicing emergency medicine physician, a computer scientist, a published author, and a federal policy advisor. This combination was rare enough that it made him something healthcare had needed for decades but rarely produced: someone who understood both what algorithms could do and what it actually felt like to stand at a patient's bedside at 2 a.m. making a decision under uncertainty.

Born in Des Moines in 1968, he earned dual degrees in Public Policy and Medical Microbiology from Stanford in 1984, then completed an MD/PhD through the joint UCSF and UC Berkeley program, followed by an emergency medicine residency at Harbor-UCLA Medical Center. He later completed a postdoctoral fellowship in medical informatics at Harvard and MIT -- the unusual combination of clinician and computer scientist that had long been theorized but rarely appeared in the same person. (Wikipedia)

He had already built something once. At age thirteen, Halamka had taught himself to program on an early personal computer. In 1981, at thirteen, he founded Ibis Research Labs, a software startup for CP/M and IBM PC platforms that he sold in 1992. The entrepreneurial instinct never left him.

He spent 20 years as Chief Information Officer of Beth Israel Deaconess Medical Center in Boston, building one of the most technologically advanced health systems in the country. He served on federal health IT advisory panels under both the Bush and Obama administrations, helping design the framework that would eventually become the HITECH Act and the national rollout of electronic health records. He held a full professorship at Harvard Medical School as International Healthcare Innovation Professor. He wrote eight books on healthcare IT. He was elected to the National Academy of Medicine in 2020. Through all of it, he continued to see patients in the emergency department.

Mayo Clinic announced his appointment in December 2019. He started January 1, 2020.

What made Halamka uniquely suited for the Platform role was not just his technical background or his credentials. It was his visceral understanding of why healthcare technology had so often failed despite promising so much. He had watched expensive systems get deployed, clinicians resist them, and promised outcomes never materialize -- not because the technology was wrong, but because the humans around it hadn't been prepared for what it would ask of them.

He had a diagnosis:

"It isn't technology; technology is good enough. It's fear of change." -- Dr. John Halamka, HealthLeaders Media, 2020

"Change management is always the hardest task." -- Dr. John Halamka, HealthLeaders Media, 2020

His prescription for Mayo Clinic Platform was therefore not primarily technical. It was organizational and relational: build a platform that convened the right people around a shared resource, with governance architecture strong enough to sustain trust across institutions, companies, patient communities, and regulators -- simultaneously. The technology, he was convinced, was the easy part.

Strategic Decision 1: The Google Partnership and the Data Enclave Model

The structural foundation of Mayo Clinic Platform had been established before Halamka arrived. In September 2019, Mayo Clinic announced a 10-year strategic partnership with Google Cloud -- a sweeping agreement that would give Mayo access to advanced cloud computing, data analytics, machine learning, and AI capabilities and make Google the cornerstone of its digital transformation. (Mayo Clinic News Network, September 2019)

The partnership accomplished two things simultaneously.

First, it built the Mayo Clinic Cloud: a Google Cloud-hosted infrastructure that centralized 1.2 million patient records in a single, secure environment that Mayo controlled entirely. Google opened a new office in Rochester, Minnesota, placing its engineers directly inside Mayo Clinic's ecosystem rather than working from a distance in Mountain View. The arrangement was deliberate: too many health IT partnerships had failed because the technology company understood algorithms but not the clinical workflow those algorithms were supposed to serve. Embedded engineers changed that dynamic.

Second, and more importantly, it created the Mayo Clinic Platform: a controlled data enclave that allowed external researchers and technology companies to work with de-identified Mayo data without

that data ever leaving Mayo's custody.

The model that emerged was called "data under glass" -- a phrase that captured both its architecture and its ethics. Third parties could bring their algorithms into the enclave. They could train their models on Mayo's patient data. But when they left, they could only take "wisdom" -- the finished, validated algorithm or the statistical insight -- not the underlying patient records. The data never moved. The enclave's "bin size of 10" de-identification standard meant that any output from the system had to be attributable to at least 10 individuals, preventing re-identification of any single patient from aggregate results. (National Academies of Sciences, 2023)

The de-identification of Mayo's data proceeded in two phases: structured data -- demographic fields, diagnosis codes, medication records -- from January through April 2020. Unstructured data -- physician notes, radiology reports, pathology findings, all the text that captured the reasoning behind decisions -- from April through August 2020. Unstructured data was harder, slower, and more consequential. It was also where the most valuable clinical insight lived.

Governance over the enclave was multilayered. A multi-stakeholder "One Table" task force reviewed all data access requests, bringing together Mayo clinical staff, legal, IT, and ethics representatives to assess each proposal. A Health Data and Technology Advisory (DaTA) Board was created in 2021 with 11 community members -- an explicit acknowledgment that data governance was not purely a technical or even a legal question, but a social one. Patients whose records were in the enclave had not consented to any specific use of their data; the community board was designed to represent their interests without making individual consent logistically impossible.

"Mayo -- an \$18 billion organization with 86,000 employees, and an urgency to do digital transformation -- has a perfect storm opportunity." -- Dr. John Halamka, HealthLeaders Media, 2020

Strategic Decision 2: Platform as Business Model -- The Convener Strategy

Infrastructure was necessary but not sufficient. Halamka understood a second-order problem clearly: building an enclave that no one used would be a very expensive filing cabinet. A platform required participants -- and participants required a compelling reason to show up.

His model was explicit on this point:

"You need a platform by which you can convene people and that may mean third-party companies, academics, or government leaders." -- Dr. John Halamka, HealthLeaders Media, 2020

"The Mayo brand is a very powerful way to bring people together." -- Dr. John Halamka, HealthLeaders Media, 2020

The "convener" strategy took concrete form in three initiatives.

Platform_Accelerate. Launched in 2022, this program embedded early-stage health AI startups inside Mayo Clinic's data ecosystem. The program began as a 20-week accelerator and later expanded to 30 weeks. In exchange for access to de-identified clinical data and expert mentorship from Mayo physicians and data scientists, startup companies validated their AI algorithms against real patient populations that they could never have accessed independently. The first cohort in early 2022 included four companies:

cliexa (cardiovascular AI for chronic disease management), Quadrant Health (EHR message triage), ScienceIO (clinical data organization tools), and Seer Medical (home-based epilepsy diagnostics). A second cohort of seven companies followed in summer 2022. By 2025, the program had accepted more than 45 startups across multiple cohorts. (Mayo Clinic News Network, T3)

Distributed Data Network. Rather than trying to aggregate all relevant health data at Mayo -- a logistically impossible and politically radioactive goal -- Halamka built a federated network of partner health systems that each contributed their own de-identified data to a shared knowledge pool. The federated architecture meant each institution retained physical and legal control of its own patient data; only the analytical insights, the trained models, and the validated findings crossed institutional boundaries. Mercy Health in St. Louis was the first member of the network.

In May 2023, the network expanded internationally, adding Hospital Israelita Albert Einstein in Brazil, Sheba Medical Center in Israel, and University Health Network in Canada -- bringing the total data pool to approximately 27 million patient lives across three continents. (Mayo Clinic News Network, May 2023) The international expansion was strategically significant: it meant that algorithms trained in the network would be validated against patient populations from diverse genetic backgrounds, healthcare systems, and disease prevalence patterns. An AI that worked only on American patients was a fundamentally limited product.

By June 2024, the distributed network had grown to nine institutional partners across four continents, representing 56 million patient lives. The solution developer ecosystem -- companies building products using Mayo Platform resources -- had nearly doubled to 81. Sixty-one healthcare providers were actively integrating platform-based solutions into clinical workflows. (Mayo Clinic News Network, June 2024)

Platform_Insights. In parallel, Mayo launched Platform_Insights, a digital innovation service allowing health systems to benchmark their quality metrics and clinical outcomes against the broader platform network. Where Platform_Accelerate served startups, Platform_Insights served established health systems that wanted to improve their own performance by learning from Mayo's data and methods.

Strategic Decision 3: The Microsoft Partnership -- AI Inside the Workflow

By 2023, Mayo Clinic Platform was running two parallel AI strategies that almost no one had named out loud. The first, built on the Google partnership, was about clinical AI: training algorithms on complex patient populations, validating them inside the data enclave, and pushing them into clinical workflows where they could reduce diagnostic error and predict disease earlier. The second, built on the Microsoft partnership, was about something entirely different: reducing the administrative friction that was consuming half of every physician's working day.

Halamka had not designed these as competing strategies. He had designed them as complementary ones -- recognizing that even the best diagnostic AI would fail clinically if the physicians who needed to use it were already drowning in documentation. You could not deploy one without the other. The Google track made AI clinically meaningful. The Microsoft track made clinical work survivable enough that clinicians could pay attention to what AI was showing them.

The problem the Microsoft partnership addressed was physician burnout.

By 2023, surveys consistently showed that American physicians spent more than half their working

hours on administrative tasks -- documentation, prior authorization, inbox management, compliance reporting -- rather than direct patient care. The documentation burden had increased dramatically with the mandated adoption of electronic health records: the same systems designed to improve information sharing had made the act of recording a clinical encounter slower and more onerous. Some estimates suggested physicians were spending two hours on administrative work for every hour with a patient.

In September 2023, Mayo Clinic became one of the first healthcare organizations to deploy Microsoft 365 Copilot, enrolling hundreds of clinical staff, doctors, and healthcare workers in Microsoft's Early Access Program. The focus was not on diagnostics -- not on making clinical decisions differently -- but on documentation: using large language models to automate form-filling, summarize patient records, draft correspondence, and reduce the friction between what a physician needed to record and how long that recording took. (Microsoft News, September 28, 2023)

Mayo's Chief Information Officer, Cris Ross, framed the initiative deliberately:

"Privacy, ethics and safety are at the forefront of Mayo Clinic's work with generative AI and large language models." -- Cris Ross, CIO, Mayo Clinic, Microsoft News, September 28, 2023

The careful framing was strategic, not bureaucratic. By 2023, the healthcare industry had watched IBM Watson Health -- once heralded as the future of AI in oncology -- fail to deliver on its clinical promises and be sold off to a private equity firm. The pattern was painfully familiar: a technology company announced revolutionary AI capabilities in healthcare, hospital systems signed contracts, and then the capabilities did not generalize from the controlled environment in which they had been demonstrated to the messier, more variable reality of actual clinical care. Mayo was determined not to repeat that pattern.

The Microsoft partnership was subsequently extended into a research collaboration with Microsoft Research and Cerebras Systems, announced in January 2025, focused on developing foundation models trained specifically on multimodal clinical data -- systems that could automatically generate chest X-ray reports, evaluate tube and line placement, and detect interval changes from prior imaging studies. The collaboration represented a qualitative shift: from deploying general-purpose AI tools adapted for healthcare to training entirely new models grounded in Mayo's clinical data from the outset.

The IBM Watson Shadow: Why History Mattered

No serious discussion of Mayo Clinic Platform's strategy was complete without understanding the failure it was designed not to replicate.

IBM Watson Health had been launched with extraordinary ambition. In 2011, IBM announced that Watson -- the same AI system that had beaten Jeopardy champions -- would be trained to recommend cancer treatment plans. Memorial Sloan Kettering Cancer Center became its primary training partner. Major hospital systems around the world signed contracts. The promise was transformative: AI that could synthesize all of medical literature and recommend optimal chemotherapy regimens for cancer patients, customized to each individual case.

By 2022, IBM had sold Watson Health to Francisco Partners, a private equity firm, at a significant loss. The failure had structural causes. Watson had been trained primarily on cases from a single institution and failed to generalize to different patient populations. Its recommendations were based on

physician-annotated training cases -- not actual patient outcomes -- meaning the system had learned to agree with oncologists at Memorial Sloan Kettering, not necessarily to produce better patient survival rates. When deployed at institutions with different patient demographics or treatment cultures, its recommendations diverged from local clinical standards in ways that clinicians could not explain or justify to patients. And the commercial pressure to announce capabilities had consistently outpaced the clinical evidence that those capabilities were real.

Halamka had spent enough time in health IT to know this history. Every structural choice in Mayo Clinic Platform -- the federated governance, the DaTA Board, the bin-size-of-10 standard, the Platform_Accelerate validation requirement, the emphasis on FDA clearance as a bar for clinical deployment -- reflected a deliberate attempt to build in the accountability mechanisms that Watson lacked.

Whether those mechanisms would hold at 56 million patient lives and 81 solution developers was a question that 2025 had not yet answered.

The change management challenge was, if anything, more acute than the technical one. Halamka had argued in 2020 that technology was the easy part -- that physician culture, institutional inertia, and the deep professional skepticism embedded in a 157-year-old medical practice were the real obstacles. By 2024, that diagnosis had been confirmed in practice: the AI tools were ready faster than the clinical workflows they were supposed to improve. Deploying an atrial fibrillation detection algorithm was straightforward. Getting cardiologists who had practiced for thirty years to treat its output as actionable -- without blindly deferring to it or reflexively overriding it -- was the work that could not be automated and was not yet done.

The AI Question: Augmentation or Replacement?

Throughout this period, Halamka returned repeatedly to a question that the broader healthcare industry was actively avoiding: what, exactly, was AI for in medicine?

His answer was consistent and direct:

"Machine learning is pattern recognition. As a clinician, I was trained to recognize patterns." -- Dr. John Halamka, HealthLeaders Media, 2020

"Machine learning is not replacing humans, but [offers] tools that augment human decision-making." -- Dr. John Halamka, HealthLeaders Media, 2020

Yet in public appearances, Halamka also deployed a more deliberately provocative formulation that generated substantial debate in medical circles:

"If your doctor can be replaced by AI, then your doctor should be replaced by AI." -- Dr. John Halamka, CXOTalk, 2023 [T3 -- content confirmed via search, full text inaccessible]

The statement, quoted out of context, sounded like a threat to the medical profession. In context, Halamka was making a narrower and more defensible argument: if a physician's primary contribution to a specific clinical decision was pattern-matching from data -- reading a chest X-ray for a standard fracture, flagging a lab value that crossed a threshold, identifying a drug interaction from a medication

list -- that specific cognitive task could and should be augmented by AI. This would free the physician to apply judgment in the places where AI fundamentally could not: the complex ambiguous presentation, the patient who needed someone to explain a diagnosis, the treatment decision that involved values as much as evidence.

Mayo Clinic's deployed clinical AI illustrated the principle in practice. The Platform had enabled algorithms for identifying patients likely to develop atrial fibrillation from EHR patterns, flagging prediabetes risk before symptoms appeared, and predicting 10-year colorectal cancer risk in patients without any prior disease history. An algorithm reducing colonoscopy polyp miss rates to a 3% error rate had been deployed -- trained by humans, validated against Mayo's patient population, and integrated into clinical workflows as a decision support tool rather than a replacement decision-maker. (CXOTalk, 2023 [T3])

In March 2024, the FDA cleared ELEFT -- the Eko Low Ejection Fraction Tool -- a stethoscope that used AI developed in conjunction with Mayo Clinic to detect low ejection fraction from cardiac sounds alone, without requiring an echocardiogram. The algorithm had been developed and validated inside Mayo's data enclave, then taken to market by Eko Health. It was precisely the model Halamka had designed: Mayo provides the data environment and the clinical validation; the ecosystem partner brings the product to market. The "wisdom" leaves the enclave. The patient data does not.

The Strategic Tensions

By 2024, Mayo Clinic Platform had achieved scale that would have seemed improbably ambitious as a five-year goal in 2020. But scale had a way of amplifying tensions that governance structures were built for smaller numbers.

Tension 1: Speed vs. Safety. The pharmaceutical and technology industries operated at fundamentally different tempos from healthcare's clinical governance. A health AI startup could iterate a model weekly. Mayo Clinic's clinical governance -- the institutional review processes, the ethics reviews, the validation requirements that had been built precisely to protect patients from unvalidated interventions -- operated on a different timeline, measured in months and sometimes years. Startups in Platform_Accelerate sometimes chafed at the pace. Mayo's clinical leadership sometimes worried that commercial pressure was pushing validation faster than the underlying evidence warranted. IBM Watson Health was the cautionary tale everyone invoked.

Tension 2: Mission vs. Revenue. Mayo Clinic was a nonprofit whose mission was patient care, not shareholder return. But Platform required sustained investment -- in infrastructure, engineering talent, governance systems, and international expansion -- and those investments needed justification within a nonprofit framework. The more commercially successful Platform became, the more it attracted scrutiny: from regulators wondering whether the data use required additional regulatory oversight, from patient advocates questioning whether patient data was being monetized in ways patients had not understood, and from within Mayo itself -- physicians and nurses who had chosen a nonprofit institution for a reason.

The financial architecture of Platform had a specific opacity problem. Mayo's public financial disclosures did not break out Platform revenue as a separate line item -- it was folded into overall institutional income. Solution developers in Platform_Accelerate paid access fees. Institutional partners contributed

data in exchange for benchmarking tools through Platform_Insights. The Google and Microsoft partnerships provided technology investment, but the commercial return to Mayo from those relationships was not publicly quantified. In the ELEFT case -- the FDA-cleared cardiac diagnostic tool developed using Mayo's enclave -- the revenue from commercializing the validated algorithm flowed to Eko Health. Mayo's return was primarily reputational, not financial, at least initially. Whether that asymmetry was sustainable for a platform requiring \$1.18 billion in annual capital expenditure was a question Mayo had not answered publicly -- and that its partners, regulators, and community board members were increasingly entitled to ask.

Tension 3: Openness vs. Competitive Advantage. The distributed data network model required deep trust from institutional partners -- trust that Mayo would not leverage the relationship to advance its own competitive position at their expense. But Mayo competed with some of those partners for patients, for research funding, for top physicians, and for the prestige that attracted all three. The "One Table" governance model was designed to manage this; whether it could sustain trust across nine institutions on four continents, with 81 solution developers and commercial partnerships with two of the world's largest technology companies, remained an open and consequential question.

Tension 4: Institutional Identity vs. Platform Logic. Halamka had been hired, in his own words, to build a "digital data business." That framing was his own: he had explicitly described Mayo Clinic Platform's goal as transforming a nonprofit hospital into a technology platform by 2025. But platform businesses operated by a logic -- grow the network, monetize interactions, extract value from ecosystem participants -- that sat in inherent tension with the values that made Mayo's brand valuable in the first place. How much of a technology company could Mayo become before its nonprofit mission, its tax status, its 157-year institutional culture, and its patients began to notice the gap?

The Current State: 2024 and Beyond

By 2024, Mayo Clinic's overall financial performance had strengthened dramatically. The 2023 operating income of \$1.1 billion represented an 82% increase over 2022 -- driven by revenue growth, improved operational efficiency, and the addition of 14,000 new staff positions. Mayo's capital expenditures reached \$1.18 billion, much of it directed toward the "Bold. Forward. Unbound." expansion of the Rochester campus -- the largest construction investment in Mayo's institutional history. (Fierce Healthcare, February 2024)

Mayo CEO Dr. Gianrico Farrugia described the moment:

"By every measure, Mayo Clinic made excellent progress in 2023 towards our goal of creating a better healthcare system." -- Dr. Gianrico Farrugia, Fierce Healthcare, February 2024

The Platform itself had grown faster than even Halamka had projected in 2020. Nine distributed data network partners. 56 million patient lives. 81 solution developers. 61 healthcare providers integrating platform solutions into their clinical workflows. At PlatforMed 2024, Mayo's annual conference for its ecosystem partners, Dr. Farrugia articulated what had become the institution's operating thesis:

"Platforms represent our greatest opportunity to bring real, lasting, and transformative change to our global healthcare system." -- Dr. Gianrico Farrugia, PlatforMed 2024 (Mayo Clinic News Network)

[T3])

Halamka was simultaneously extending Platform's frontier. The Microsoft Research and Cerebras Systems collaboration announced in January 2025 moved the ambition decisively from deploying existing AI tools to training entirely new foundation models grounded in Mayo's clinical data from the start -- models that could reshape radiology, pathology, and genomics in ways that off-the-shelf large language models could not.

That announcement landed five years after Halamka had arrived in Rochester with a question about whether a 157-year-old medical institution could remake itself as a technology platform. By early 2025, the platform existed, the data network spanned four continents, and the foundation model collaboration meant Mayo was now attempting to build AI that no one had built before -- trained not on internet text but on decades of real clinical complexity, at a scale no single institution had previously assembled.

The governance board still met. The DaTA community representatives still reviewed access requests. The bin-size-of-10 rule still applied. But the platform Halamka was now governing was roughly ten times the size of the one he had designed those rules for in 2020. The institutional trust that the "data under glass" architecture was built to protect had never been worth more -- or been more exposed to the kinds of failures, commercial pressures, and governance gaps that had ended IBM Watson Health's ambitions at far smaller scale.

The question was not whether the platform was succeeding. The question was whether the structures Halamka had built five years earlier, when the answer was still theoretical, were built for what it had become.

Discussion Questions

1. Dr. Halamka described Mayo Clinic Platform's goal as transforming Mayo into a "digital data business." Is this description compatible with Mayo's nonprofit mission? Where does the line between mission-driven innovation and commercialization fall, and who should draw it?
2. The "data under glass" model was designed to prevent patient data from leaving Mayo's custody. What are the governance risks of this model as it scales to 56 million patient lives across 9 institutions on 4 continents? Who should govern it -- and who should NOT?
3. Halamka argued that "change management is always the hardest task" -- not the technology. What specific change management challenges would you anticipate in deploying AI tools to physicians at a 157-year-old institution? How would you address them differently than technology-first approaches have?
4. IBM Watson Health was once heralded as the future of AI in oncology. It failed and was divested in 2022. What structural factors differentiate Mayo Clinic Platform's AI strategy from Watson's approach? Are those factors sufficient, or are there risks the case does not address?
5. Halamka said: "If your doctor can be replaced by AI, then your doctor should be replaced by AI." Evaluate this claim. What is the strongest argument in its favor? The strongest argument against? What does your answer imply for how healthcare systems should be designed?
6. Mayo Clinic competes with some of its distributed data network partners for patients, talent, and

research funding. How should Mayo manage this conflict of interest? What specific governance mechanisms would you recommend, and what should their limits be?

7. Mayo's 2023 capital expenditure was \$1.18 billion, much of it directed to physical campus expansion. How should Mayo balance physical infrastructure investment against digital platform investment? Is there a tension between the two strategic directions, or are they mutually reinforcing?

8. Platform_Accelerate offers startups access to Mayo's patient data and clinical brand in exchange for validation work. Who benefits more from this arrangement -- Mayo or the startups? Is this exchange equitable? Should startups receive equity, revenue sharing, or other compensation for the value their validated products add to the Mayo ecosystem?

Case prepared by: Akshay Kumar | ITEC-617 | Kogod School of Business, American University | Spring 2026

AI assistance: Claude Code (claude-sonnet-4-6) with humanizer and wshobson-agents data-storytelling skills
Research date: April 26-27, 2026 | Source verification: See sources/Source_Registry.md