Mayo Clinic: AI Platform Transformation in Healthcare

Two-Page Submission Report · ITEC-617 Individual Assignment

Student: Akshay Kumar | Course: ITEC-617 | Kogod School of Business, American University
Date: May 3, 2026 | Repository: github.com/akbknight/mayo-clinic-dt-case-study

Case Study Quality

The case follows Dr. John Halamka, who became the inaugural President of Mayo Clinic Platform on January 1, 2020. His mandate: turn Mayo's patient data -- accumulated over 157 years of complex, rare-case medicine -- into the foundation of a technology platform that external companies could use to develop and validate clinical AI. He arrived three months before COVID-19 hit.

What makes the case interesting is not the technology. It is the structural problem: Mayo's data has clinical depth because its nonprofit mission generated it, but monetizing that data through platform logic requires commercial relationships that sit in direct tension with that mission. Halamka had to build governance that sustained trust with patients, regulators, competing health systems, and two of the world's largest technology companies -- simultaneously, with no precedent for how to do it.

The three strategic decisions -- the Google data enclave, the distributed network expansion, and the Microsoft generative AI deployment -- are not parallel choices. They represent a deliberate two-track strategy: Google built the infrastructure for clinical diagnostic AI; Microsoft addressed the workflow problem that made clinical AI adoption impossible if physicians were too buried in documentation to pay attention to what it showed them. Halamka was explicit that both tracks were necessary.

The unresolved tensions are real, not rhetorical: speed vs. safety (startup iteration cycles vs. clinical governance timelines), mission vs. revenue (nonprofit identity vs. commercial platform logic), and openness vs. competitive advantage (requiring trust from health systems that also compete with Mayo). IBM Watson Health runs throughout the case not as a rhetorical device but as a structural comparison -- Watson failed for specific, diagnosable reasons, and every governance choice in Mayo's design is a deliberate attempt not to repeat them. The case leaves the protagonist governing, in 2025, a platform ten times the size of the one he designed the governance rules for in 2020.

Quality Verification

I verified every specific figure against its source before treating it as fact.

Confirmed: The $17.9B revenue, 10.2% growth, and $1.1B operating income are consistent across multiple independent sources and trace to Fierce Healthcare's February 2024 report on Mayo's 2023 annual disclosure. The IBM Watson MD Anderson $62M project discontinuation is documented in multiple public sources including STAT News -- not a fabrication. Dr. Halamka's direct quotes came from a full content fetch of the HealthLeaders interview; they are verbatim and dated.

Caught and flagged: Mayo Clinic's own website returned 403 errors on every automated fetch attempt. This created a cluster of T3 sources -- confirmed to exist through search but not fully accessible at text. The "replace your doctor" quote and the Platform_Accelerate cumulative startup figures (45+ by 2025) are T3: confirmed by search snippet but hedged accordingly in the document. The AI market size projection in the supplement ($22.45B ? $208.23B by 2030) has no inline citation; I flagged this gap in the source registry rather than fabricating a source.

Most significant correction: The first draft framed Mayo's Platform as a clear institutional success. I pushed back on that framing. The revision surfaces specific unresolved tensions -- the revenue model opacity, the governance scaling problem, the change management gap between AI tool readiness and physician workflow adoption -- rather than endorsing the strategy as proven.

Reflection & Learning

Claude compressed days of library research into hours. But everything it produced required human curation -- deciding what to trust, what to cut, and which tensions mattered for a business audience. It drafted well; it did not judge. The hardest editorial work -- deciding what to cut, how to frame unresolved questions, how much to hedge T3 claims -- required human judgment every time.

The most unexpected finding was the irony: Mayo Clinic News Network blocked every automated fetch with 403 errors. Their own website protects information from AI scrapers in the same way their platform protects patient data from unauthorized access. The "data under glass" principle applies to their public relations apparatus too.

On Mayo Clinic's DT: The real innovation is governance, not technology. Mayo designed oversight structures that made commercial AI partnerships legitimate to patients, regulators, and physicians simultaneously, without a regulatory playbook for how to do it. Halamka's line -- "change management is always the hardest task" -- describes not just physician adoption of AI tools but the entire governance challenge of the platform itself.

What I would do differently: download PDFs of paywalled academic sources (MIT Sloan, NEJM, JAMA) before starting. University library access would have moved four T3 sources to T1. Starting verification before drafting -- not after -- would also have reduced editorial rework.

AI Tool Usage & Process

Three tools, three roles.

Claude Code handled research analysis, narrative drafting, and verification. The most effective prompt structure: "Extract only direct quotes and data points from this source. Flag anything you are inferring rather than reading directly." This produced cleaner drafts than open-ended writing prompts. When I asked Claude to "write the case," the first output used standard AI writing patterns -- significance inflation, promotional framing, vague attributions. Running the draft through the humanizer skill (which identifies 29 specific AI writing patterns) removed them. I ran it three times on the main case before the prose read cleanly.

Perplexity was fastest for targeted fact-checking. "Is the IBM Watson MD Anderson $62M figure documented in public reporting?" Confirmed with citations in seconds -- more efficient than Claude for quick verification questions with a known answer.

GitHub Copilot handled repository mechanics -- YAML configuration, markdown formatting, CHANGELOG entries. Nothing intellectually demanding, but it removed tedious setup friction.

The most useful process change I made: treating verification as a parallel track rather than a final step. Every claim flagged at time of writing (T1/T2/T3) rather than reviewed as a batch afterward. It slowed drafting slightly but caught attribution gaps before they were embedded in finished prose.